

Welcome to Machine Learning!

COMP 4630 | Winter 2025


Charlotte Curtis

What is this course about?

- Continuing the supervised/unsupervised learning algorithms from COMP 3652, with a focus on **Neural Networks**
- First half: the history, theory, and math behind neural networks
- Second half: applications of NNs in computer vision, natural language processing, and more

This is not (just) a course on building models using libraries like TensorFlow or PyTorch, it is a course on understanding the theory

How did I get involved with ML?



What do you want to learn about ML?



Grade Assessment

| Component | Weight |
|-----------------------|---------|
| Assignments | 3 x 10% |
| Midterm (theory) exam | 20% |
| Journal club | 15% |
| Final project | 35% |

Bonus marks may be awarded for *substantial* corrections to materials, submitted as pull requests

Course materials repo: <https://github.com/mru-comp4630/w25>

Textbooks and other readings

Primary Textbook:

- [Hands on Machine Learning with Scikit-Learn and Tensorflow](#)
- [Associated GitHub repo](#)

More mathy details:

- [Deep Learning](#)

Journal club list:

- [MRU Library Reading List](#)

Machine Learning Project Checklist

[Appendix A of the hands-on textbook](#)

1. **Frame the problem** and look at the big picture.
2. **Get the data.**
3. **Explore the data** to gain insights.
4. **Prepare the data** to better expose the underlying data patterns to Machine Learning algorithms.
5. Explore many different models and short-list the best ones.
6. Fine-tune your models and combine them into a great solution.
7. Present your solution.
8. Launch, monitor, and maintain your system.

1. Look at the big picture

Example Dataset: California housing prices (1990)

? Discussion questions:

- How does the company expect to use and benefit from this model?
- What is the current solution?
- What kind of ML task is this?
- What kind of performance measure should we use?

2. Get the data

For this class, we'll use readily available datasets. Some sources are:

- [UCI Machine Learning Repository](#)
- [Kaggle](#)
- [Google Dataset Search](#)
- Various Government open data portals (e.g. [Calgary](#), [Alberta](#), [Canada](#))

After fetching the data, set aside a test set and **don't look at it**.

"Get the data" can often be a huge task in itself!

2a. Set aside a test set

? Discussion questions:

- Why do we need an independent test set?
 - Avoid data snooping bias
 - [Relevant XKCD](#)
- Why would we use a random seed?
- What is naive about simply selecting a random sample?
- What else could we do?
- What is stratified sampling?

Side tangent: Sampling bias

- Simple example: assume 80% of population likes cilantro
- Goal: ensure our sample is representative of the population, $\pm 5\%$

The [binomial distribution](#) can be used to model the probability of choosing k people who like cilantro from n total participants:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}, \text{ where } \binom{n}{k} = \frac{n!}{k!(n - k)!}$$

Side tangent: Sampling bias continued

$P(X = k)$ is the probability mass function, and the corresponding cumulative distribution function is just the sum up to k :

$$P(X \leq k) = \sum_{i=0}^k \binom{n}{i} p^i (1-p)^{n-i}$$

By adding together $P(X \leq 0.8 - 0.05)$ and $P(X \geq 0.8 + 0.05)$, we get the probability of sampling bias (by the $\pm 5\%$ definition).

This is also my excuse to review some probability theory and notation

3. Explore the data

? Discussion questions:

- What do you notice about the data?
- Do the values make sense for the labels?
- Is the scale of the features comparable? Does this matter?
- What possible biases might be present in the data?

3a. Look for correlations

The **Pearson correlation coefficient** is a measure of the linear correlation between two variables X and Y (commonly denoted as r):

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

where \bar{x} and \bar{y} are the sample means of X and Y , respectively.

- What do correlations of 0, 1, and -1 mean?
- What are some limitations of Pearson correlation?

4. Prepare the data

General goals:

- Handle missing data, and maybe outliers
- Drop irrelevant features
- Combine features using domain knowledge
- Apply various transformations (e.g. scaling, encoding)
- Apply scaling when necessary

4a. Handling missing data

In the book 3 options are listed to handle the NaN values:

```
housing.dropna(subset=["total_bedrooms"], inplace=True) # option 1
housing.drop("total_bedrooms", axis=1)                  # option 2
median = housing["total_bedrooms"].median()             # option 3
housing["total_bedrooms"].fillna(median, inplace=True)
```

? Discussion questions:

- What is each option doing?
- What are the pros and cons of each option?
- Which one should we choose?

4b. Handling non-numeric data

Most of the math in ML algorithms is based on numbers, so we need to convert text and categorical attributes to numbers. This is called **encoding**.

? Discussion questions:

- Which columns of our data are categorical?
- What methods could we use to convert them to numbers?
- What are the assumptions about the various encoding methods?

4c. Scaling the data

Many ML algorithms don't like features with vastly different scales. Common scaling methods are **min-max scaling** and **standardization**.

*Important: scaling is **computed** on the training set and **applied** to the validation and test sets - they are not scaled independently!*

? Discussions questions:

- What are the bounds of each method?
- Which method is more affected by outliers?
- How would you decide which method to use?

4e. Standardization details

A general Gaussian distribution is given by:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

where μ is the mean and σ is the standard deviation. The standard normal distribution is a special case where $\mu = 0$ and $\sigma = 1$, reducing the equation to:

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}$$

4f. Other transformations

- **Log transformation:** useful for data that is heavily skewed
- Also **square root, squaring, etc.:** try to remove heavy tails
- **Feature engineering:** combining features to create new ones
- **Binning:** turning continuous data into discrete categories
 - Possibly using K-means clustering
 - Relies on domain knowledge
- Best to create a **transformation pipeline** and apply it to the data rather than saving the transformed data