

# Training Models with Regression and Gradient Descent

---

COMP 4630 | Winter 2025

Charlotte Curtis

# Overview

---

- Linear Regression and the Normal Equation
- Gradient Descent and its various flavours
- References and suggested reading:
  - [Scikit-learn book](#):
    - Chapter 4: Training Models
  - [Deep Learning Book](#)
    - Section 5.1.4: Linear Regression

# Linear Regression

---

Unlike most models, linear regression has a **closed-form** solution called the **Normal Equation**:

$$\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where

- $\hat{\theta}$  are the weights of the model minimizing the **cost function**
- $\mathbf{y}$  is the vector of target values
- $\mathbf{X}$  is the **design matrix** of feature values

*As usual, different sources use different notation, e.g.  $\mathbf{w}$  or  $\phi$  instead of  $\theta$ .*

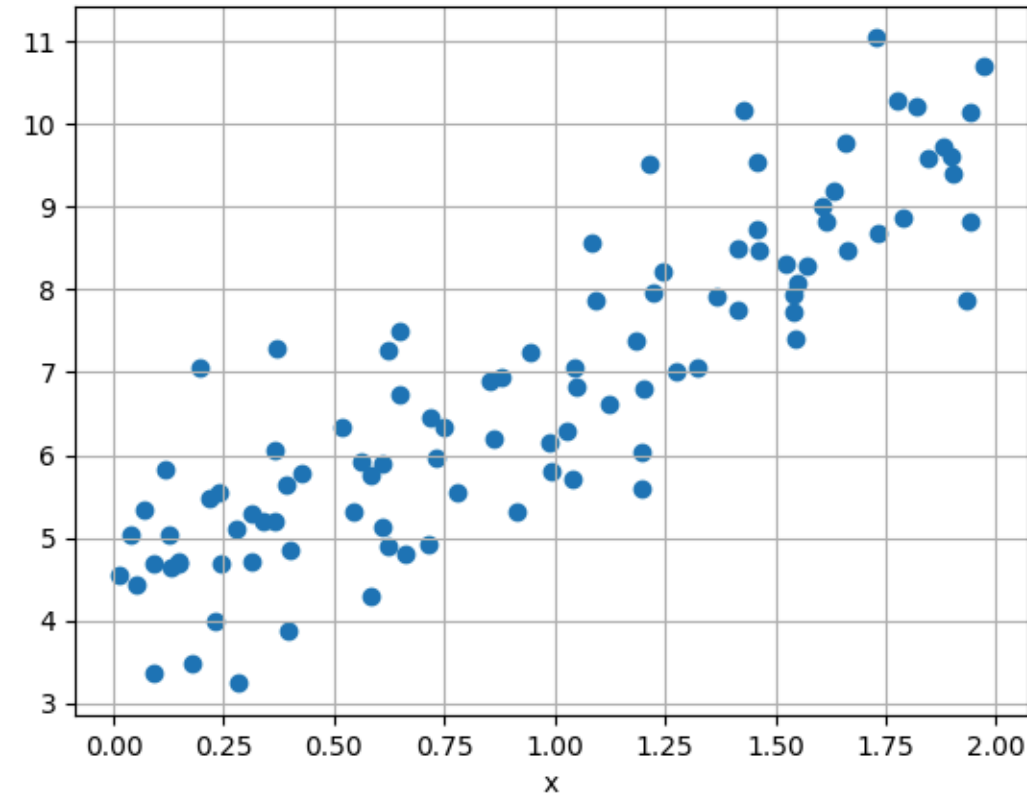
Consider the 1-d case:

$$\hat{y} = \theta_0 + \theta_1 x$$

we want the values of  $\theta_0$  and  $\theta_1$  that minimize the **Mean Square Error** between the actual and predicted  $y$  values:

$$MSE = \frac{1}{m} \sum_{i=1}^m (\hat{y} - y_i)^2$$

$$MSE = \frac{1}{m} \sum_{i=1}^m (\theta_0 + \theta_1 x_i - y_i)^2$$



# Solving for $\theta_0$ and $\theta_1$




# Solving for $\theta_0$ and $\theta_1$

After some algebraic gymnastics, we get:

$$\theta_1 = \frac{\mu_y \sum_m x_i - \sum_m x_i y_i}{\mu_x \sum_m x_i - \sum_m x_i^2}$$

$$\theta_0 = \mu_y - \theta_1 \mu_x$$

where  $\mu_x$  and  $\mu_y$  are the means of the  $x$  and  $y$  values, respectively.



# Expanding to matrix form

---

Instead of the scalar  $x$  or even vector  $\mathbf{x}$ , we can use a **design matrix**  $\mathbf{X}$  to represent the feature values:

$$\mathbf{X} = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1n} \\ x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m1} & x_{m2} & \cdots & x_{mn} \end{bmatrix}$$

where each row is an instance (sample) and each column is a feature.

*It is common for the first column to be all ones, representing the bias term*

# Back to the linear regression problem...

---

- We can rewrite the estimate in matrix notation:

$$\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta}$$

- The MSE can be written as:

$$MSE = \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{m} (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

where we've used the trick of substituting  $\mathbf{a}^T \mathbf{a} = \sum_i a_i^2$

-  Find the gradient of the MSE w.r.t  $\boldsymbol{\theta}$ , set it to zero, and solve for  $\boldsymbol{\theta}$



# Properties of matrices and their transpose

---

The following properties are useful for solving linear algebra problems:

- $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$
- $(\mathbf{A} + \mathbf{B})^T = \mathbf{A}^T + \mathbf{B}^T$
- $(\mathbf{A}^{-1})^T = (\mathbf{A}^T)^{-1}$
- $(\mathbf{A}^T)^T = \mathbf{A}$

Additionally, any matrix or vector multiplied by  $\mathbf{I}$  is unchanged.

# The Normal Equation

---

We made it! The **Normal Equation** is again:

$$\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- No optimization is required to find the optimal  $\theta$
- Limitations:
  - $\mathbf{X}^T \mathbf{X}$  must be invertible and small enough to fit in memory
  - The computational complexity is (at least)  $O(n^3)$
- Even in linear regression problems, it is common to use **gradient descent** instead due to these limitations

# Gradient Descent

---

The goal of gradient descent is still to minimize the cost function, but it follows an iterative process:

1. Start with a random  $\theta$
2. Calculate the gradient  $\nabla_{\theta}$  for the current  $\theta$
3. Update  $\theta$  as  $\theta = \theta - \eta \nabla_{\theta}$
4. Repeat 2-3 until some stopping criterion is met

where  $\eta$  is the **learning rate**, or the size of step to take in the direction opposite the gradient.

# Stochastic Gradient Descent

---

- Standard or **batch** gradient descent uses the entire training set to calculate the gradient for each instance at every step
- **Stochastic Gradient Descent** uses a single random instance at each step:
  - i. Start with a random  $\theta$
  - ii. Pick a random instance  $\mathbf{x}_i$  (row in the design matrix)
  - iii. Calculate the gradient  $\nabla_{\theta}$  for the current  $\theta$  and  $\mathbf{x}_i$
  - iv. Update  $\theta$  as  $\theta = \theta - \eta \nabla_{\theta}$
  - v. Repeat 2-4 until some stopping criterion is met

# Mini-batch Gradient Descent

- **Mini-batch** gradient descent uses a random *subset* of the training set
- Less chaotic than stochastic, but faster than batch
- Most common type of gradient descent used in practice



# Gradient Descent Hyperparameters

---

- The **learning rate**  $\eta$  - size of step taken
- No rule that it needs to be constant! A simple **learning schedule** is to decrease  $\eta$  over time, e.g.:

$$\eta = \frac{t_0}{t + t_1}$$

where  $t$  is the current iteration and  $t_0$  and  $t_1$  are hyper-parameters

- For mini-batch, the **batch size** is another hyper-parameter
- The number of **epochs**, or times to process the entire training set

# Stopping Criteria

- The simplest stopping criterion is to set a maximum number of epochs
- **Early stopping** is another option:
  - Evaluate on a validation set at regular intervals
  - Stop when the validation error starts to increase
- The comparison between training and validation performance can also help prevent **overfitting**



# Loss functions

- The **loss function** is the function being minimized by gradient descent
- MSE is **convex** and guaranteed to have a single global minimum, but many other loss functions have multiple local minima
- The relative scale of the features can affect the convergence:




Figure 4-7. Gradient Descent with (left) and without (right) feature scaling



# Higher-order Polynomials

---

- Higher order polynomials can be solved with the Normal Equation as well:

$$y = \theta_0 + \theta_1 x + \theta_2 x^2 + \dots + \theta_n x^n$$

- Just include the higher order terms in  $\mathbf{X}$
- This is still a linear regression problem because the coefficients are linear!
- Risk of **overfitting** the data
- Easy way to regularize: drop one or more of the higher order terms

# Regularization

---

- If the model fits the training data *too* well, but doesn't generalize to new data, it is **overfitting**
- **Regularization** imposes additional constraints on the weights
- Example: **Ridge Regression** adds a term to the loss function:

$$J(\theta) = MSE(\theta) + \alpha \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

where  $\alpha$  is the **regularization parameter**

- The regularization term is only added during training, not evaluation

*Note: the term **cost** function is often used instead of **loss** function*

# Logistic regression and beyond

---

Logistic regression is a binary classifier that uses the **logistic function** (aka **sigmoid function**) to map the output to a range of 0 to 1:

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

resulting in the loss function (trust me on this one):

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)]$$

where  $\hat{p}_i = \sigma(\theta^T \mathbf{x}_i)$  is the probability that instance  $i$  is positive.

The gradient of the loss function for logistic regression is:

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m (\sigma(\theta^T \mathbf{x}_i) - y_i) \mathbf{x}_i$$

- There is no (known) analytical solution this time, but we can still use gradient descent!
- In this case it's still convex, so we don't have to worry about local minima
- In general, for a loss function to work with gradient descent, it must be:
  - **Continuous** and
  - **Differentiable**
  - ... at the locations where you evaluate it

# **Next up: Backpropagation!**

---