

# Classification loss functions and metrics

---

COMP 4630 | Winter 2025

Charlotte Curtis

# Overview

---

- All the derivation thus far has been for mean squared error
- **Cross-entropy loss** is more appropriate for classification problems
- References and suggested reading:
  - [Scikit-learn book](#): Chapter 4, training models
  - [Scikit-learn docs](#): Log loss
  - [Deep Learning Book](#): Sections 3.1, 3.8, and 6.2

# Statistics review: Expected value

---

The **expected value** of some function  $f(x)$  when  $x$  is distributed as  $P(x)$  is given in discrete form as:

$$\mathbb{E}[f(x)] = \sum_x P(x) f(x)$$

where the sum is over all possible values of  $x$ .

In continuous form, this is an integral:

$$\mathbb{E}[f(x)] = \int p(x) f(x) dx$$

# Binary case: Bernoulli distribution

---

- If a random variable  $x$  has a  $p$  probability of being 1 and a  $1 - p$  probability of being 0, then  $x$  is distributed as a **Bernoulli distribution**:

$$P(x) = p^x (1 - p)^{1-x} = \begin{cases} p & \text{for } x = 1 \\ 1 - p & \text{for } x = 0 \end{cases}$$

- The expected value of  $x$  is then:

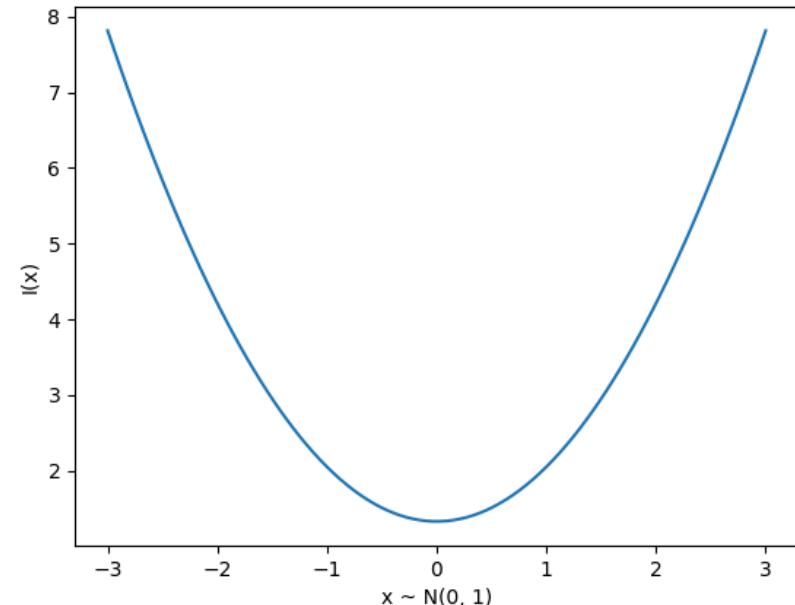
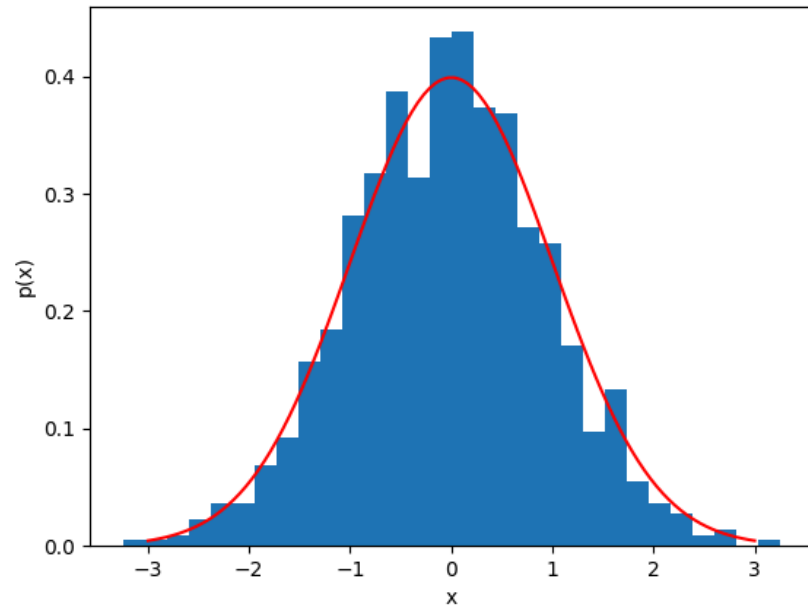
$$\mathbb{E}[x] = 0 \cdot (1 - p) + 1 \cdot p = p$$

# Information theory

---

Originally developed for message communication, with the intuition that less likely events carry more **information**, defined for a single event as:

$$I(x) = -\log P(x)$$



# Entropy

- We can measure the **expected information** of a distribution  $P(x)$  as:

$$H(X) = \mathbb{E}[I(x)] = -\mathbb{E}_{x \sim P}[\log P(x)]$$

- This is called the **Shannon entropy**
- Measured in bits (base 2) or nats (base  $e$ )
- 🎲 Find the entropy of a bernoulli distribution



# Cross-entropy

---

- The **KL divergence** is a measure of the *extra* information needed to encode a message from a true distribution  $P(x)$  using an approximate distribution  $Q(x)$ :

$$D_{KL}(P||Q) = \mathbb{E}_{x \sim P} \left[ \log \frac{P(x)}{Q(x)} \right] = \mathbb{E}_{x \sim P} [\log P(x) - \log Q(x)]$$

- The **cross-entropy** is a simplification that drops the term  $\log P(x)$ :

$$H(P, Q) = -\mathbb{E}_{x \sim P} [\log Q(x)]$$

- Minimizing the cross-entropy is equivalent to minimizing the KL divergence
- If  $P(x) = Q(x)$ , then  $D_{KL}(P||Q) = 0$  and  $H(P, Q) = H(P)$


# Cross-entropy loss

For a true label  $y \in \{0, 1\}$  and predicted  $\hat{p} \in [0, 1]$ , the cross-entropy loss is:

$$\begin{aligned}\mathcal{L}(y, \hat{p}) &= -\mathbb{E}_y[\log P(x)] \\ &= -y \log \hat{p} - (1 - y) \log(1 - \hat{p})\end{aligned}$$

where  $\hat{p} = \sigma(\mathbf{w}^T \mathbf{h} + b)$  is the output of the final layer of a neural network (thresholded to obtain the prediction  $\hat{y}$ )

*This is also called **log loss** or **binary cross-entropy***





# Multiclass case

---

- For  $K$  classes, the output is a vector  $\hat{\mathbf{p}}$  with  $\hat{p}_i = P(y = i|\mathbf{x})$
- The cross-entropy loss is then:

$$\mathcal{L}(\mathbf{y}, \hat{\mathbf{p}}) = - \sum_{i=1}^K y_i \log \hat{p}_i$$

- For a one-hot encoded vector  $\mathbf{y}$ , this simplifies to:

$$\mathcal{L}(\mathbf{y}, \hat{\mathbf{p}}) = - \log \hat{p}_k$$

where  $k$  is the index of the true class

# The softmax function

---

- For binary classification, the sigmoid function  $\sigma(z) = \frac{1}{1+e^{-z}}$  is used to predict the probability of the positive class
- For multiclass classification, the **softmax function** is used:

$$\hat{p}_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

where  $z_i = \mathbf{w}_i^T \mathbf{h} + b_i$  is the output of neuron  $i$  in the final layer before the activation function is applied

- This means that  $K$  neurons are needed in the final layer, one for each class

# Terminology for evaluation

---

- **True positive:** predicted positive, label was positive ( $TP$ ) ✓
- **True negative:** predicted negative, label was negative ( $TN$ ) ✓
- **False positive:** predicted positive, label was negative ( $FP$ ) ✗ (type I)
- **False negative:** predicted negative, label was positive ( $FN$ ) ✗ (type II)
- **Accuracy** is the fraction of correct predictions, given as:

$$\text{accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

# Precision and recall

---

- **Precision:** Out of all the positive **predictions**, how many were correct?

$$\text{precision} = \frac{TP}{TP + FP}$$

- **Recall:** Out of all the positive **labels**, how many were correct?

$$\text{recall} = \frac{TP}{TP + FN}$$

- **Specificity:** Out of all the negative **labels**, how many were correct?

$$\text{specificity} = \frac{TN}{TN + FP}$$

# Confusion matrix

---

	Predicted Positive	Predicted Negative
True Positive	TP	FN
True Negative	FP	TN

- The axes might be reversed, but a good predictor will have strong diagonals
- There's also the **F1 score**, or harmonic mean of precision and recall:

$$F1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

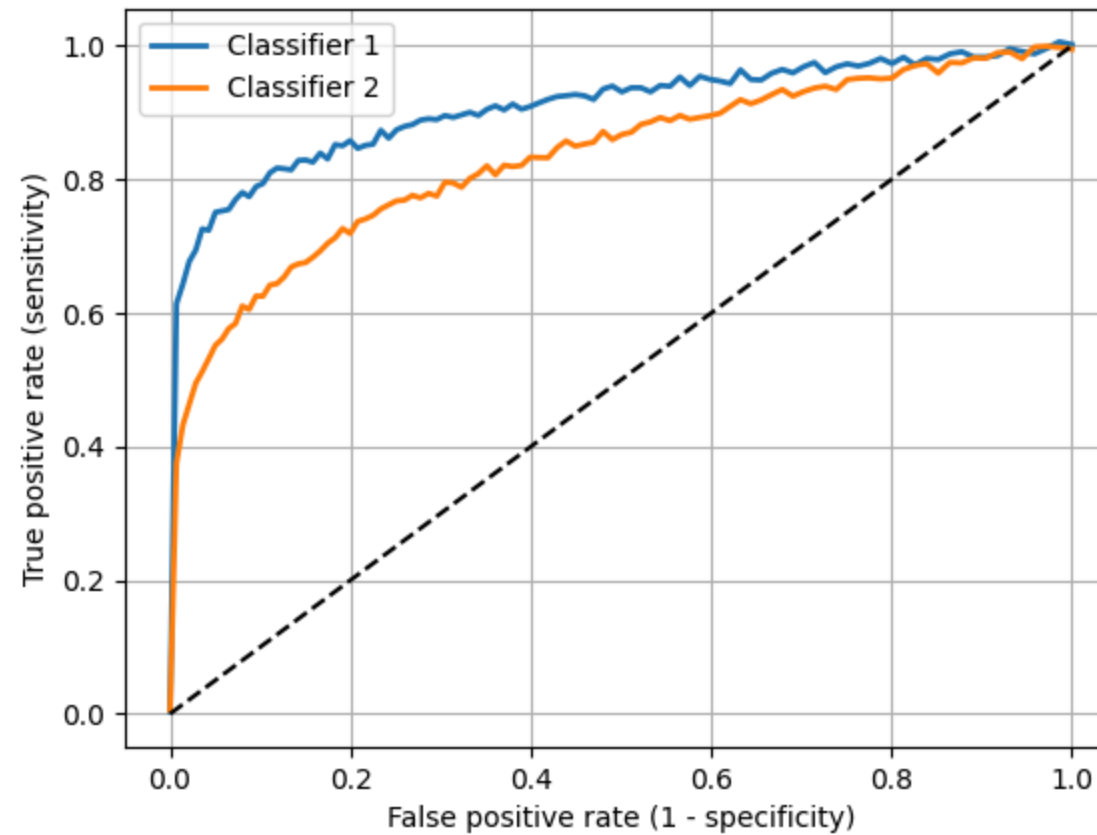
# ROC Curves

---

- The **receiver operating characteristic** curve is a plot of the **true positive rate** (recall or sensitivity) vs. **false positive rate** (1 - specificity) as the detection threshold changes
- The diagonal is the same as random guessing
- A perfect classifier would hug the top left corner

*Fun fact: the name comes from WWII radar operators, where true positives were airplanes and false positives were noise*

# Which classifier is better?



**Next up: Convolution and NN frameworks**

---