

Generative AI, ethics, and policies

COMP 4630 | Winter 2025

Charlotte Curtis

Overview

- Discussion on AI usage, ethics, and policies
- Autoencoders
- Generative adversarial networks (GANs)
- Diffusion models
- References and suggested reading:
 - [Scikit-learn book](#): Chapter 17
 - [Deep Learning book](#): Chapter 14, 20

Guiding questions for discussion

- ? What are some ways that AI can be helpful? Harmful?
- ? What are some ways that AI is used?
 - Axes: intent (malicious <-> innocent), impact (harmful <-> helpful)
- ? What concerns do you have about AI in general?
 - What about in education, specifically at MRU?
- ? What would you like to see in an AI policy at MRU?

A case study on impact: Vibe coding

Costs* of:

- Voice to text
- Code generation
- Debugging
- at ~3 Wh/request

*Carbon footprint, \$\$, scalability, security, etc



Andrej Karpathy 
@karpathy



There's a new kind of coding I call "vibe coding", where you fully give in to the vibes, embrace exponentials, and forget that the code even exists. It's possible because the LLMs (e.g. Cursor Composer w Sonnet) are getting too good. Also I just talk to Composer with SuperWhisper so I barely even touch the keyboard. I ask for the dumbest things like "decrease the padding on the sidebar by half" because I'm too lazy to find it. I "Accept All" always, I don't read the diffs anymore. When I get error messages I just copy paste them in with no comment, usually that fixes it. The code grows beyond my usual comprehension, I'd have to really read through it for a while. Sometimes the LLMs can't fix a bug so I just work around it or ask for random changes until it goes away. It's not too bad for throwaway weekend projects, but still quite amusing. I'm building a project or webapp, but it's not really coding - I just see stuff, say stuff, run stuff, and copy paste stuff, and it mostly works.

4:17 PM · Feb 2, 2025 · **4.5M** Views

Autoencoders

- Like encoder-decoder networks, but the input and output are the same, minimizing:

$$\mathcal{L}(\mathbf{x}, f(g(\mathbf{x})))$$

where f is the encoder and g is the decoder

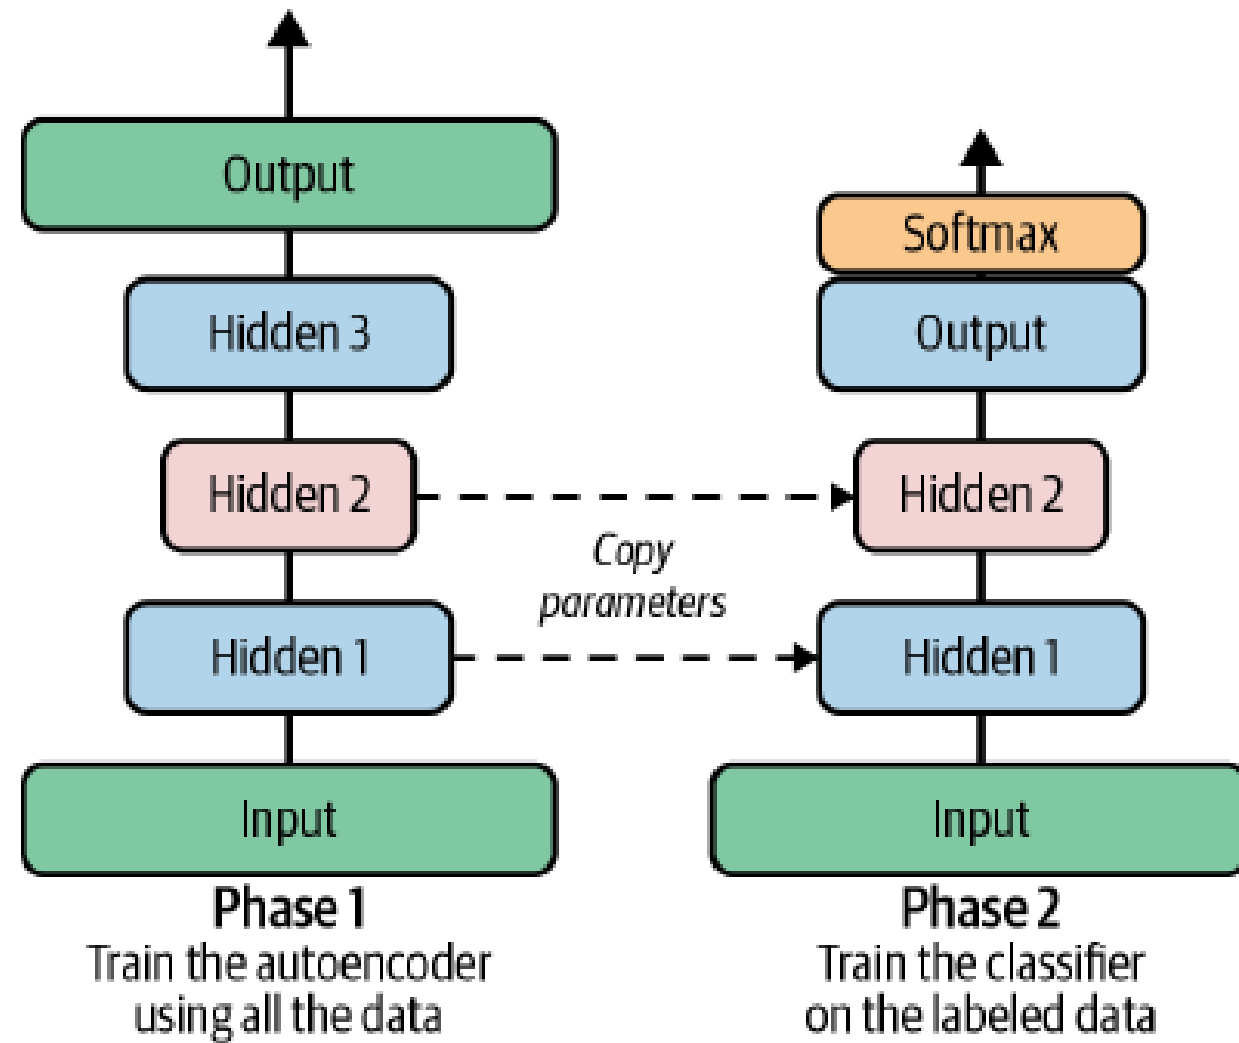
- Autoencoders have been around [since the 80s](#)
- **?** If we're just training a model to copy its input, what's the point?

What can we do with autoencoders?

- Autoencoders learn a **latent space** representation of the data, provided the latent space is **lower dimension** than the input (undercomplete)
- **Regularization** can also help the model learn a useful representation
- Once trained, the model can be used for:
 - Data compression
 - Information retrieval
 - Anomaly detection
 - Unsupervised pretraining

Unsupervised pretraining

- Rather than training (e.g.) a classifier from scratch, we can first train an **autoencoder**, then use the **encoder** part as the first half of the classifier
- ? Why might this be useful?
- ? What other applications can you think of?



More autoencoder tricks

- **Sparse autoencoders** introduce a sparsity penalty to the loss function:

$$\mathcal{L}(\mathbf{x}, g(f(\mathbf{x}))) + \Omega(\mathbf{h})$$

where $\mathbf{h} = f(\mathbf{x})$ and Ω is a sparsity penalty

- The model can now learn useful features without constraining latent dimensions
- **Weight tying** takes advantage of the symmetric nature of autoencoders by forcing the weights of the decoder to be the same as the encoder
- **Denoising autoencoders** can be created by adding noise to the input:

$$\mathcal{L}(\mathbf{x}, g(f(\tilde{\mathbf{x}})))$$

- Autoencoders can also be CNNs, RNNs, or just about any other kind of network!

Variational Autoencoders

- Encoder produces a **distribution** (μ, σ)
- Decoder samples from $N(\mu, \sigma)$ to produce the output
- Loss includes a **KL divergence** term to push the distribution towards normal




Figure 17-11. A variational autoencoder (left) and an instance going through it (right)

Generative Adversarial Networks (GANs)

- VAEs can be used as generators (e.g. images), so how can we train them to generate what we actually want?
- In [2014](#), Ian Goodfellow et. al proposed training two networks simultaneously: a **generator** and a **discriminator**
 - The generator tries to produce realistic looking images
 - The discriminator is trained to classify images as real or fake
- Two-phase training loop:
 - Equal numbers of real and fake images used to train discriminator
 - Discriminator weights frozen while generator produces images with label of "real" - backpropagation updates only generator weights

GAN training challenges

- Training is a **zero-sum** game, where each "player" receives a payoff that the discriminator tries to maximize and the generator tries to minimize
- Equilibrium is reached when the discriminator performance is 50%, but **oscillations** can occur and convergence is not guaranteed
- **Mode collapse** happens when the generator figures out that it can produce a single image that fools the discriminator
- Still an active area of research!

"And since many factors affect these complex dynamics, GANs are very sensitive to the hyperparameters: you may have to spend a lot of effort fine-tuning them. In fact, that's why I used RMSProp rather than Nadam when compiling the models: when using Nadam, I ran into a severe mode collapse." - Aurélien Géron

Diffusion Models

- First formalized in 2015 by Jascha Sohl-Dickstein et. al
- Not popular until 2020 when Jonathan Ho et. al introduced **DDPM**
- Concept: take an image and gradually add noise until it's unrecognizable, then train a model to reverse the process

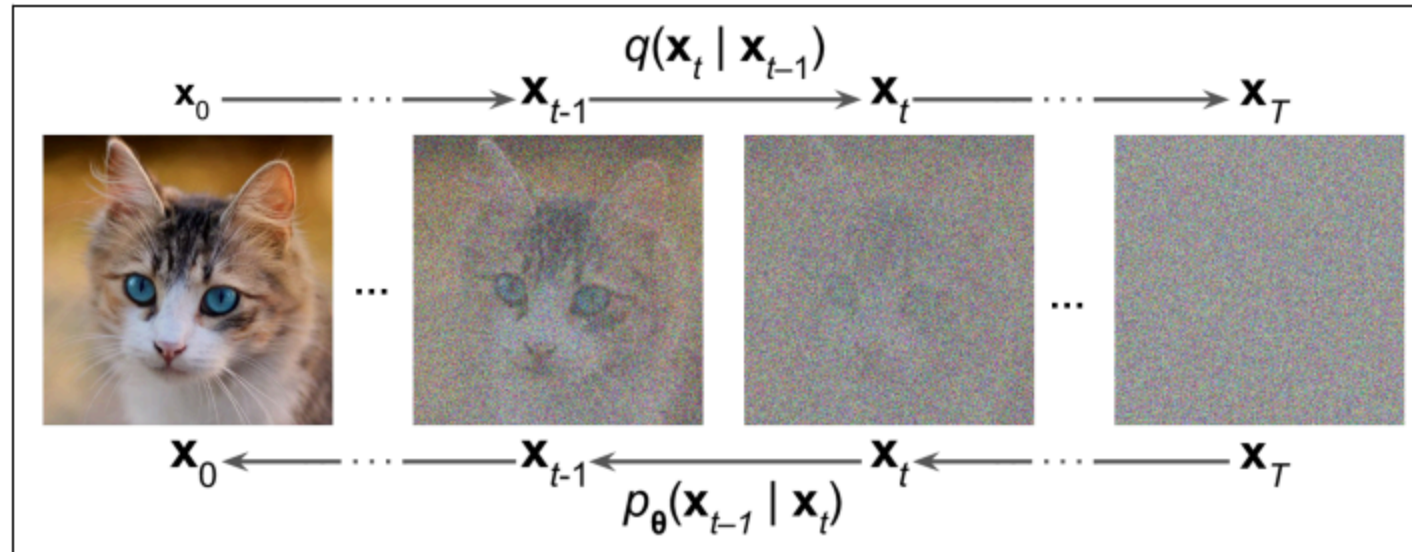


Figure 17-20. The forward process q and reverse process p

CLIP guidance

- DDPM, GANs, and VAEs are all capable of generating images given a **class label** that is part of the training data
- Image generation + NLP = magic
- Contrastive Language-Image Pretraining **CLIP** was trained on 400 million image-caption pairs to predict which caption goes with which image
- Final result is a general purpose model that can go both ways
- Actually **open sourced** by OpenAI!

That's all, next up is your project presentations!
